

```
In [177...]  
import warnings  
warnings.filterwarnings("ignore")  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns
```

```
In [180...]  
adult=pd.read_csv(r"D:\adult.data.csv",na_values='?',skipinitialspace=True)  
adult
```

Out[180...]

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	United-States
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
48837	39	Private	215419	Bachelors	13	Divorced	Prof-specialty	Not-in-family	White	Female	0	0	36	United-States
48838	64	Nan	321403	HS-grad	9	Widowed	Nan	Other-relative	Black	Male	0	0	40	United-States
48839	38	Private	374983	Bachelors	13	Married-civ-spouse	Prof-specialty	Husband	White	Male	0	0	50	United-States
48840	44	Private	83891	Bachelors	13	Divorced	Adm-clerical	Own-child	Asian-Pac-Islander	Male	5455	0	40	United-States
48841	35	Self-emp-inc	182148	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	60	United-States

48842 rows × 15 columns

```
In [182...]  
adult.head()
```

Out[182...]

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba

```
In [184...]  
adult.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48842 entries, 0 to 48841
Data columns (total 15 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   age         48842 non-null   int64  
 1   workclass    46043 non-null   object  
 2   fnlwgt       48842 non-null   int64  
 3   education    48842 non-null   object  
 4   education-num 48842 non-null   int64  
 5   marital-status 48842 non-null   object  
 6   occupation   46033 non-null   object  
 7   relationship  48842 non-null   object  
 8   race         48842 non-null   object  
 9   sex          48842 non-null   object  
 10  capital-gain 48842 non-null   int64  
 11  capital-loss 48842 non-null   int64  
 12  hours-per-week 48842 non-null   int64  
 13  native-country 47985 non-null   object  
 14  income        48842 non-null   object  
dtypes: int64(6), object(9)
memory usage: 5.6+ MB
```

```
In [186]: adult.shape
```

```
Out[186]: (48842, 15)
```

```
In [188]: adult.dtypes
```

```
Out[188]: age            int64
workclass        object
fnlwgt           int64
education        object
education-num    int64
marital-status   object
occupation       object
relationship     object
race             object
sex              object
capital-gain    int64
capital-loss    int64
hours-per-week   int64
native-country   object
income           object
dtype: object
```

```
In [190]: adult.isna().sum()
```

```
Out[190]: age            0
workclass        2799
fnlwgt           0
education        0
education-num    0
marital-status   0
occupation       2809
relationship     0
race             0
sex              0
capital-gain    0
capital-loss    0
hours-per-week   0
native-country   857
income           0
dtype: int64
```

```
In [192]: adult_new=adult.fillna(method='ffill')
adult_new
```

Out[192...]

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	United-States
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
48837	39	Private	215419	Bachelors	13	Divorced	Prof-specialty	Not-in-family	White	Female	0	0	36	United-States
48838	64	Private	321403	HS-grad	9	Widowed	Prof-specialty	Other-relative	Black	Male	0	0	40	United-States
48839	38	Private	374983	Bachelors	13	Married-civ-spouse	Prof-specialty	Husband	White	Male	0	0	50	United-States
48840	44	Private	83891	Bachelors	13	Divorced	Adm-clerical	Own-child	Asian-Pac-Islander	Male	5455	0	40	United-States
48841	35	Self-emp-inc	182148	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	60	United-States

48842 rows × 15 columns



In [194...]

adult\_new.columns

Out[194...]

```
Index(['age', 'workclass', 'fnlwgt', 'education', 'education-num',
       'marital-status', 'occupation', 'relationship', 'race', 'sex',
       'capital-gain', 'capital-loss', 'hours-per-week', 'native-country',
       'income'],
      dtype='object')
```

In [196...]

```
adult_new=adult_new.drop(columns=['fnlwgt','relationship','native-country'],axis=1)
adult_new
```

Out[196...]

	age	workclass	education	education-num	marital-status	occupation	race	sex	capital-gain	capital-loss	hours-per-week	income
<b>0</b>	39	State-gov	Bachelors	13	Never-married	Adm-clerical	White	Male	2174	0	40	<=50K
<b>1</b>	50	Self-emp-not-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	White	Male	0	0	13	<=50K
<b>2</b>	38	Private	HS-grad	9	Divorced	Handlers-cleaners	White	Male	0	0	40	<=50K
<b>3</b>	53	Private	11th	7	Married-civ-spouse	Handlers-cleaners	Black	Male	0	0	40	<=50K
<b>4</b>	28	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	Black	Female	0	0	40	<=50K
...	...	...	...	...	...	...	...	...	...	...	...	...
<b>48837</b>	39	Private	Bachelors	13	Divorced	Prof-specialty	White	Female	0	0	36	<=50K.
<b>48838</b>	64	Private	HS-grad	9	Widowed	Prof-specialty	Black	Male	0	0	40	<=50K.
<b>48839</b>	38	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	White	Male	0	0	50	<=50K.
<b>48840</b>	44	Private	Bachelors	13	Divorced	Adm-clerical	Asian-Pac-Islander	Male	5455	0	40	<=50K.
<b>48841</b>	35	Self-emp-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	White	Male	0	0	60	>50K.

48842 rows × 12 columns

In [198...]

```
adult_new.describe()
```

Out[198...]

	age	education-num	capital-gain	capital-loss	hours-per-week
<b>count</b>	48842.000000	48842.000000	48842.000000	48842.000000	48842.000000
<b>mean</b>	38.643585	10.078089	1079.067626	87.502314	40.422382
<b>std</b>	13.710510	2.570973	7452.019058	403.004552	12.391444
<b>min</b>	17.000000	1.000000	0.000000	0.000000	1.000000
<b>25%</b>	28.000000	9.000000	0.000000	0.000000	40.000000
<b>50%</b>	37.000000	10.000000	0.000000	0.000000	40.000000
<b>75%</b>	48.000000	12.000000	0.000000	0.000000	45.000000
<b>max</b>	90.000000	16.000000	99999.000000	4356.000000	99.000000

In [200...]

```
adult_new.isna().sum()
```

Out[200...]

age	0
workclass	0
education	0
education-num	0
marital-status	0
occupation	0
race	0
sex	0
capital-gain	0
capital-loss	0
hours-per-week	0
income	0
dtype: int64	

In [202...]

```
adult_new['income']=adult_new['income'].str.strip().str.replace('.',' ', regex=False)
adult_new
```

Out[202...]

	age	workclass	education	education-num	marital-status	occupation	race	sex	capital-gain	capital-loss	hours-per-week	income
<b>0</b>	39	State-gov	Bachelors	13	Never-married	Adm-clerical	White	Male	2174	0	40	<=50K
<b>1</b>	50	Self-emp-not-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	White	Male	0	0	13	<=50K
<b>2</b>	38	Private	HS-grad	9	Divorced	Handlers-cleaners	White	Male	0	0	40	<=50K
<b>3</b>	53	Private	11th	7	Married-civ-spouse	Handlers-cleaners	Black	Male	0	0	40	<=50K
<b>4</b>	28	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	Black	Female	0	0	40	<=50K
...	...	...	...	...	...	...	...	...	...	...	...	...
<b>48837</b>	39	Private	Bachelors	13	Divorced	Prof-specialty	White	Female	0	0	36	<=50K
<b>48838</b>	64	Private	HS-grad	9	Widowed	Prof-specialty	Black	Male	0	0	40	<=50K
<b>48839</b>	38	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	White	Male	0	0	50	<=50K
<b>48840</b>	44	Private	Bachelors	13	Divorced	Adm-clerical	Asian-Pac-Islander	Male	5455	0	40	<=50K
<b>48841</b>	35	Self-emp-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	White	Male	0	0	60	>50K

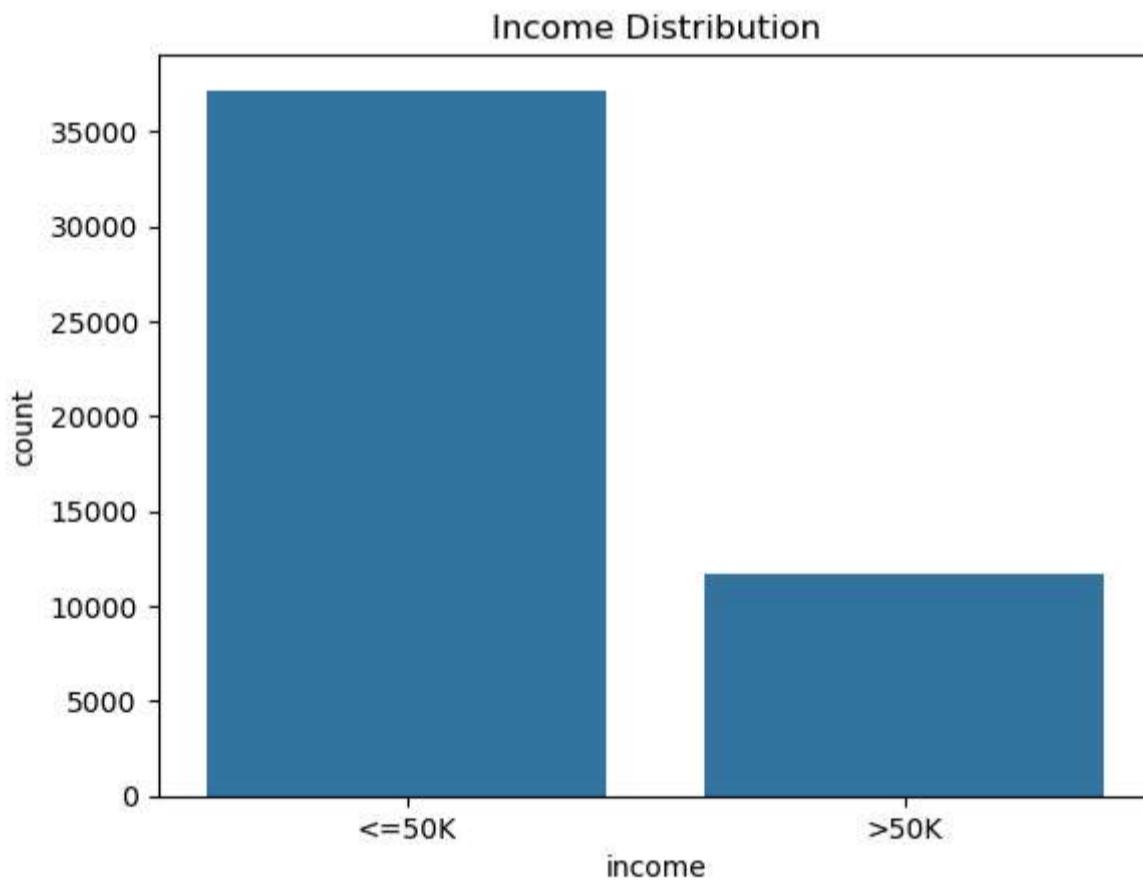
48842 rows × 12 columns

In [204...]

```
sns.countplot(x='income', data=adult_new)
plt.title('Income Distribution')
```

Out[204...]

Text(0.5, 1.0, 'Income Distribution')

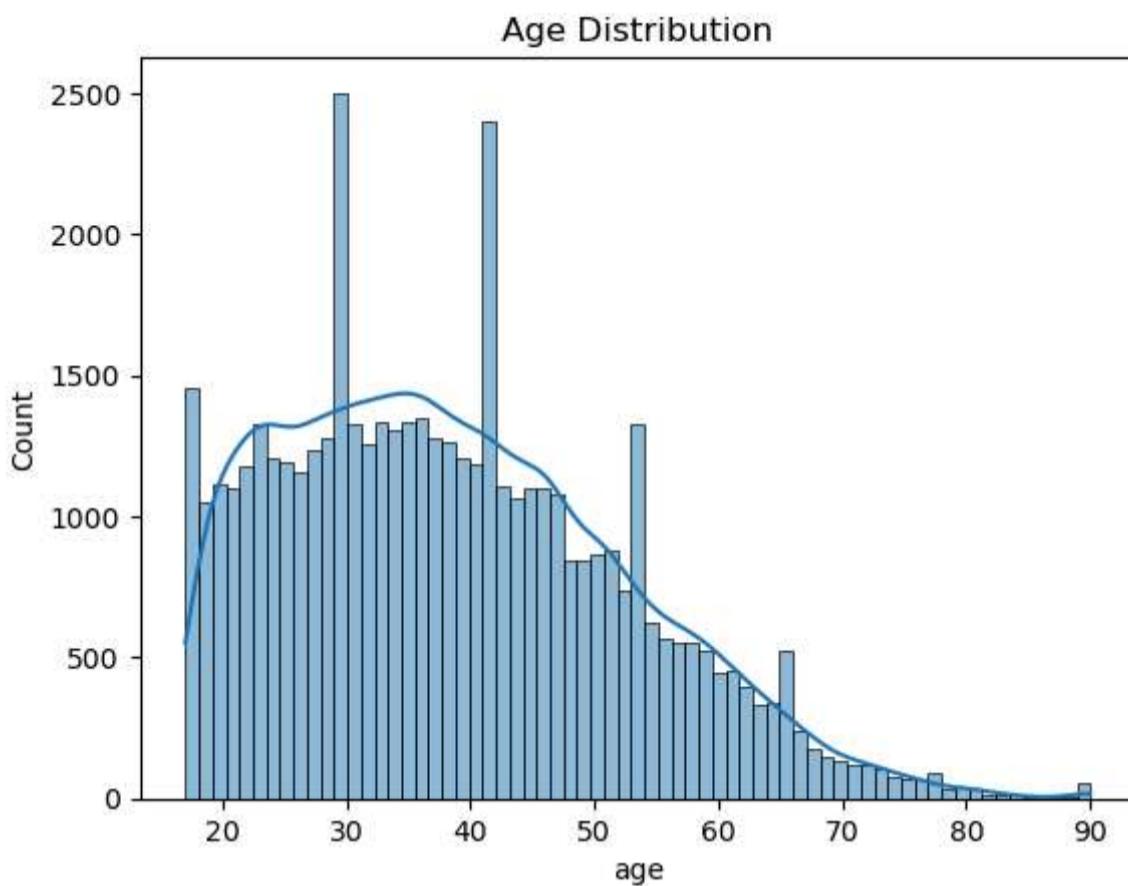


In [205...]

```
sns.histplot(adult_new['age'], kde=True)
plt.title('Age Distribution')
```

Out[205...]

Text(0.5, 1.0, 'Age Distribution')



```
In [210]: from sklearn.preprocessing import LabelEncoder
```

```
In [212]: encode=LabelEncoder()
adult_new['income']=encode.fit_transform(adult_new['income'])
adult_new['income']
```

```
Out[212]:
0      0
1      0
2      0
3      0
4      0
..
48837  0
48838  0
48839  0
48840  0
48841  1
Name: income, Length: 48842, dtype: int32
```

```
In [214]: adult_new.columns
```

```
Out[214]:
Index(['age', 'workclass', 'education', 'education-num', 'marital-status',
       'occupation', 'race', 'sex', 'capital-gain', 'capital-loss',
       'hours-per-week', 'income'],
      dtype='object')
```

```
In [216]: adult_new=pd.get_dummies(adult_new,drop_first=True)
adult_new
```

	age	education-num	capital-gain	capital-loss	hours-per-week	income	workclass_Local-gov	workclass_Never-worked	workclass_Private	workclass_Self-emp-inc	...	occu
0	39	13	2174	0	40	0	False	False	False	False	False	...
1	50	13	0	0	13	0	False	False	False	False	False	...
2	38	9	0	0	40	0	False	False	True	False	False	...
3	53	7	0	0	40	0	False	False	True	False	False	...
4	28	13	0	0	40	0	False	False	True	False	False	...
..	...	...	...	...	...	...	...	...	...	...	...	...
48837	39	13	0	0	36	0	False	False	True	False	False	...
48838	64	9	0	0	40	0	False	False	True	False	False	...
48839	38	13	0	0	50	0	False	False	True	False	False	...
48840	44	13	5455	0	40	0	False	False	True	False	False	...
48841	35	13	0	0	60	1	False	False	False	True	False	...

48842 rows × 52 columns

```
In [228]: x=adult_new.drop(columns=['income'])
x
```

Out[228...]

	age	education-num	capital-gain	capital-loss	hours-per-week	workclass_Local-gov	workclass_Never-worked	workclass_Private	workclass_Self-emp-inc	workclass_Self-emp-not-inc
0	39	13	2174	0	40	False	False	False	False	False
1	50	13	0	0	13	False	False	False	False	True
2	38	9	0	0	40	False	False	True	False	False
3	53	7	0	0	40	False	False	True	False	False
4	28	13	0	0	40	False	False	True	False	False
...	...	...	...	...	...	...	...	...	...	...
48837	39	13	0	0	36	False	False	True	False	False
48838	64	9	0	0	40	False	False	True	False	False
48839	38	13	0	0	50	False	False	True	False	False
48840	44	13	5455	0	40	False	False	True	False	False
48841	35	13	0	0	60	False	False	False	True	False

48842 rows × 51 columns

In [230...]

```
y=adult_new['income']
y
```

Out[230...]

```
0      0
1      0
2      0
3      0
4      0
 ..
48837  0
48838  0
48839  0
48840  0
48841  1
Name: income, Length: 48842, dtype: int32
```

In [258...]

```
from sklearn.model_selection import train_test_split
```

In [260...]

```
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.8,random_state=42,stratify=y)
```

In [262...]

```
x_test.head()
```

Out[262...]

	age	education-num	capital-gain	capital-loss	hours-per-week	workclass_Local-gov	workclass_Never-worked	workclass_Private	workclass_Self-emp-inc	workclass_Self-emp-not-inc
40421	30	9	0	0	40	False	False	True	False	False
47738	54	12	0	0	39	False	False	False	False	False
518	21	10	0	0	35	False	False	True	False	False
8564	35	9	2885	0	40	False	False	True	False	False
31355	42	10	0	0	45	True	False	False	False	False

5 rows × 51 columns

In [264...]

```
x_train.shape
```

Out[264...]

(39073, 51)

In [266...]

```
x_test.shape
```

Out[266...]

(9769, 51)

In [269...]

```
from sklearn.preprocessing import StandardScaler
```

In [271...]

```
scaler=StandardScaler()
scaler
```

Out[271...]

```
StandardScaler(i ?)
StandardScaler()
```

```
In [273... scaled_x_train=scaler.fit_transform(x_train)
scaled_x_train

Out[273... array([[-0.12009008,  1.13821044, -0.14407503, ..., -0.09327335,
       0.4115339 , -1.42415255],
       [ 1.26824482, -0.80402454, -0.14407503, ..., -0.09327335,
       0.4115339 ,  0.70217197],
       [ 1.04903405, -0.41557754, -0.14407503, ..., -0.09327335,
       0.4115339 ,  0.70217197],
       ...,
       [ 1.63359611, -0.02713055, -0.14407503, ..., -0.09327335,
       0.4115339 ,  0.70217197],
       [-0.04701982,  0.36131645, -0.14407503, ..., -0.09327335,
       0.4115339 , -1.42415255],
       [ 0.31833147, -0.02713055, -0.14407503, ..., -0.09327335,
       -2.42993349,  0.70217197]])
```

```
In [275... scaled_x_test=scaler.fit_transform(x_test)
scaled_x_test

Out[275... array([[-0.62594162, -0.43451805, -0.14860513, ..., -0.08434099,
       0.41257776,  0.71243367],
       [ 1.11198537,  0.73872873, -0.14860513, ..., -0.08434099,
       0.41257776,  0.71243367],
       [-1.27766425, -0.04343579, -0.14860513, ..., -0.08434099,
       0.41257776, -1.40363945],
       ...,
       [ 1.4740535 , -0.43451805, -0.14860513, ..., -0.08434099,
       0.41257776,  0.71243367],
       [-0.69835525, -0.04343579, -0.14860513, ..., -0.08434099,
       0.41257776, -1.40363945],
       [ 1.4740535 , -0.04343579, -0.14860513, ..., -0.08434099,
       -2.42378549,  0.71243367]])
```

## Logistic Regression

```
In [279... from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix,classification_report

In [281... model=LogisticRegression(solver='liblinear')
model

Out[281... ▾ LogisticRegression ⓘ ⓘ
LogisticRegression(solver='liblinear')
```

```
In [283... model.fit(x_train,y_train)

Out[283... ▾ LogisticRegression ⓘ ⓘ
LogisticRegression(solver='liblinear')
```

```
In [285... y_pred=model.predict(x_test)
y_pred
```

```
Out[285... array([0, 0, 0, ..., 0, 0, 0])
```

```
In [289... accuracy_log=accuracy_score(y_test,y_pred)
accuracy_log
```

```
Out[289... 0.8497287337496161
```

```
In [291... print("Accuracy :",accuracy_log)
```

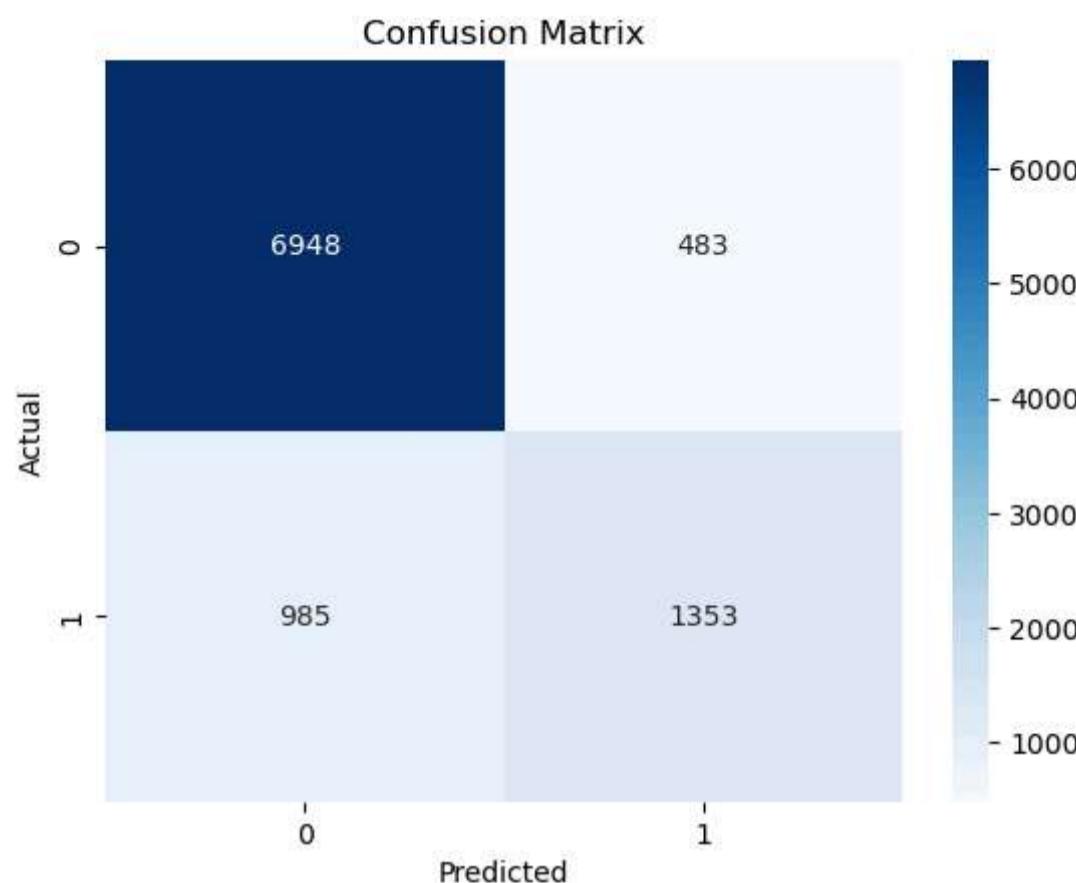
```
Accuracy : 0.8497287337496161
```

```
In [294... confusion_matrix(y_test,y_pred)
```

```
Out[294... array([[6948,  483],
       [ 985, 1353]], dtype=int64)
```

```
In [296... sns.heatmap(confusion_matrix(y_test,y_pred),annot=True,fmt='d',cmap='Blues')
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
```

```
Out[296... Text(50.72222222222214, 0.5, 'Actual')
```



```
In [297... print("classification Report :\n")
print(classification_report(y_test,y_pred))
```

classification Report :

	precision	recall	f1-score	support
0	0.88	0.94	0.90	7431
1	0.74	0.58	0.65	2338
accuracy			0.85	9769
macro avg	0.81	0.76	0.78	9769
weighted avg	0.84	0.85	0.84	9769

## Random Forest

```
In [301... from sklearn.ensemble import RandomForestClassifier
```

```
In [303... model=RandomForestClassifier(n_estimators=50,random_state=100)
model
```

```
Out[303... RandomForestClassifier
RandomForestClassifier(n_estimators=50, random_state=100)
```

```
In [305... model.fit(x_train,y_train)
```

```
Out[305... RandomForestClassifier
RandomForestClassifier(n_estimators=50, random_state=100)
```

```
In [306... y_pred=model.predict(x_test)
y_pred
```

```
Out[306... array([0, 0, 0, ..., 1, 0, 0])
```

```
In [310... accuracy_rand=accuracy_score(y_test,y_pred)
accuracy_rand
print("Accuracy :",accuracy_rand)
```

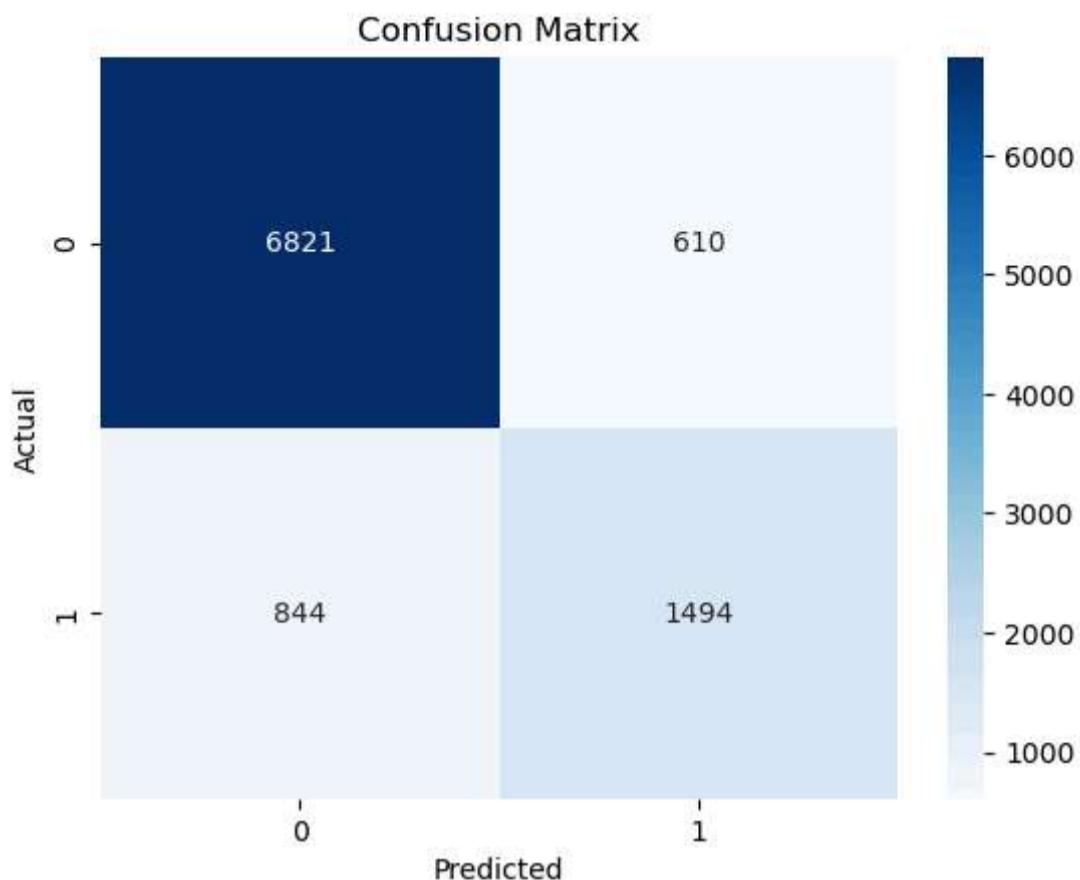
Accuracy : 0.8511618384686253

```
In [313... confusion_matrix(y_test,y_pred)
```

```
Out[313... array([[6821,  610],
 [ 844, 1494]], dtype=int64)
```

```
In [315... sns.heatmap(confusion_matrix(y_test,y_pred),annot=True,fmt='d',cmap='Blues')
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
```

```
Out[315... Text(50.72222222222214, 0.5, 'Actual')
```



```
In [317]: print("classification Report :\n")
print(classification_report(y_test,y_pred))
```

classification Report :

	precision	recall	f1-score	support
0	0.89	0.92	0.90	7431
1	0.71	0.64	0.67	2338
accuracy			0.85	9769
macro avg	0.80	0.78	0.79	9769
weighted avg	0.85	0.85	0.85	9769

## Decision tree classification

```
In [320]: from sklearn.tree import DecisionTreeClassifier
```

```
In [322]: model=DecisionTreeClassifier(criterion='gini',max_depth=3)
model
```

```
Out[322]: 
▼ DecisionTreeClassifier ⓘ ⓘ
DecisionTreeClassifier(max_depth=3)
```

```
In [324]: model.fit(x_train,y_train)
```

```
Out[324]: 
▼ DecisionTreeClassifier ⓘ ⓘ
DecisionTreeClassifier(max_depth=3)
```

```
In [326]: y_pred=model.predict(x_test)
y_pred
```

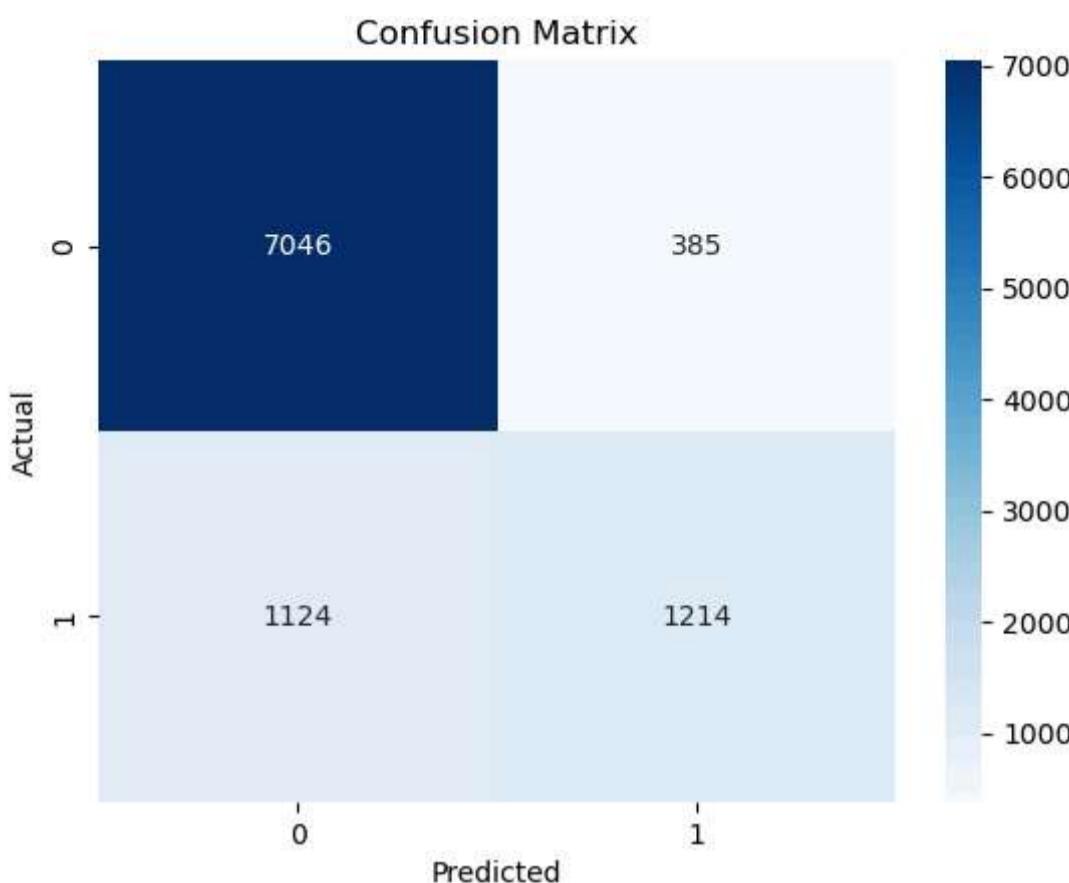
```
Out[326]: array([0, 0, 0, ..., 0, 0, 0])
```

```
In [329]: confusion_matrix(y_test,y_pred)
```

```
Out[329]: array([[7046, 385],
 [1124, 1214]], dtype=int64)
```

```
In [331]: sns.heatmap(confusion_matrix(y_test,y_pred),annot=True,fmt='d',cmap='Blues')
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
```

```
Out[331]: Text(50.72222222222214, 0.5, 'Actual')
```



```
In [334]: accuracy_tree=accuracy_score(y_test,y_pred)
accuracy_tree
```

```
Out[334]: 0.8455317842153751
```

```
In [336]: print("classification Report :\n")
print(classification_report(y_test,y_pred))
```

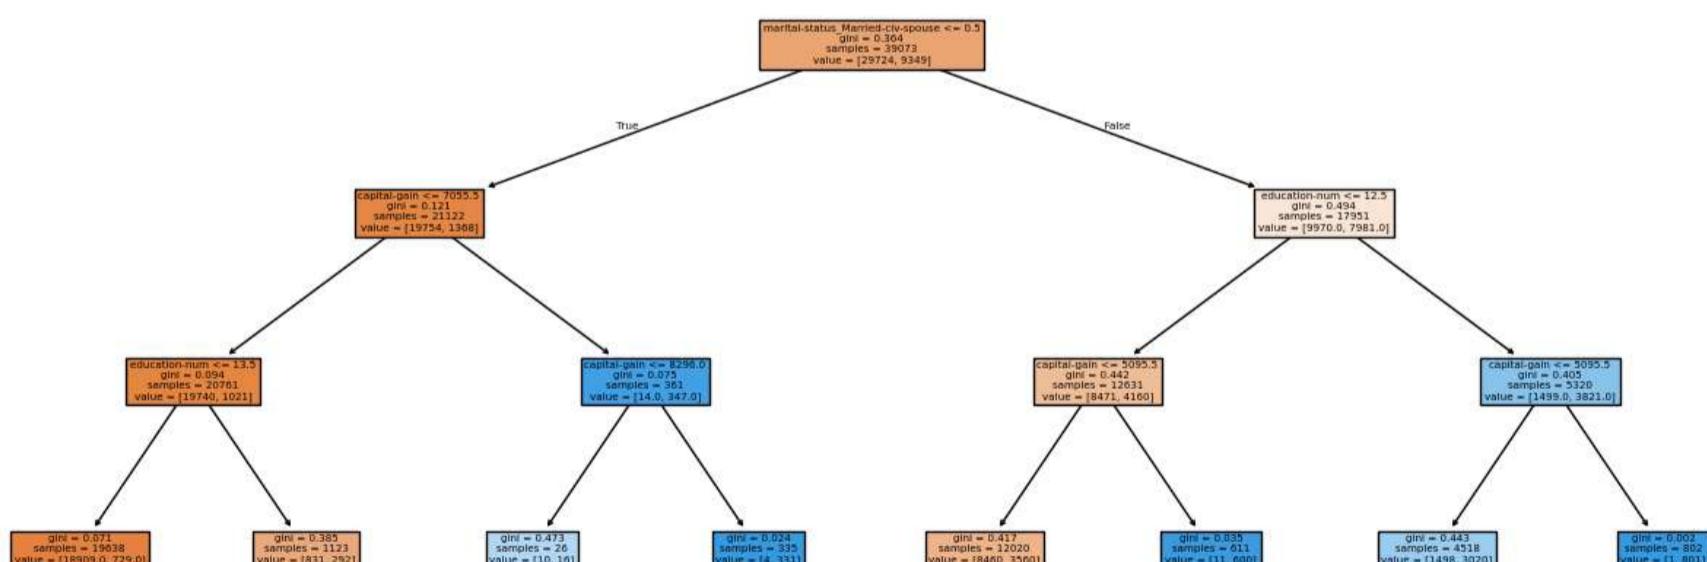
classification Report :

	precision	recall	f1-score	support
0	0.86	0.95	0.90	7431
1	0.76	0.52	0.62	2338
accuracy			0.85	9769
macro avg	0.81	0.73	0.76	9769
weighted avg	0.84	0.85	0.83	9769

### Tree plotting

```
In [339]: from sklearn.tree import plot_tree
```

```
In [341]: plt.figure(figsize=(16,6))
plot_tree(model,feature_names=x.columns,filled=True)
plt.show()
```



### K-Nearest Neighbors (KNN)

```
In [344]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [346]: model=KNeighborsClassifier(n_neighbors=3)
model
```

```
Out[346... ▾ KNeighborsClassifier ⓘ ⓘ
KNeighborsClassifier(n_neighbors=3)
```

```
In [348... model.fit (scaled_x_train,y_train)
```

```
Out[348... ▾ KNeighborsClassifier ⓘ ⓘ
KNeighborsClassifier(n_neighbors=3)
```

```
In [350... y_pred=model.predict(scaled_x_test)
y_pred
```

```
Out[350... array([0, 0, 0, ..., 0, 0, 0])
```

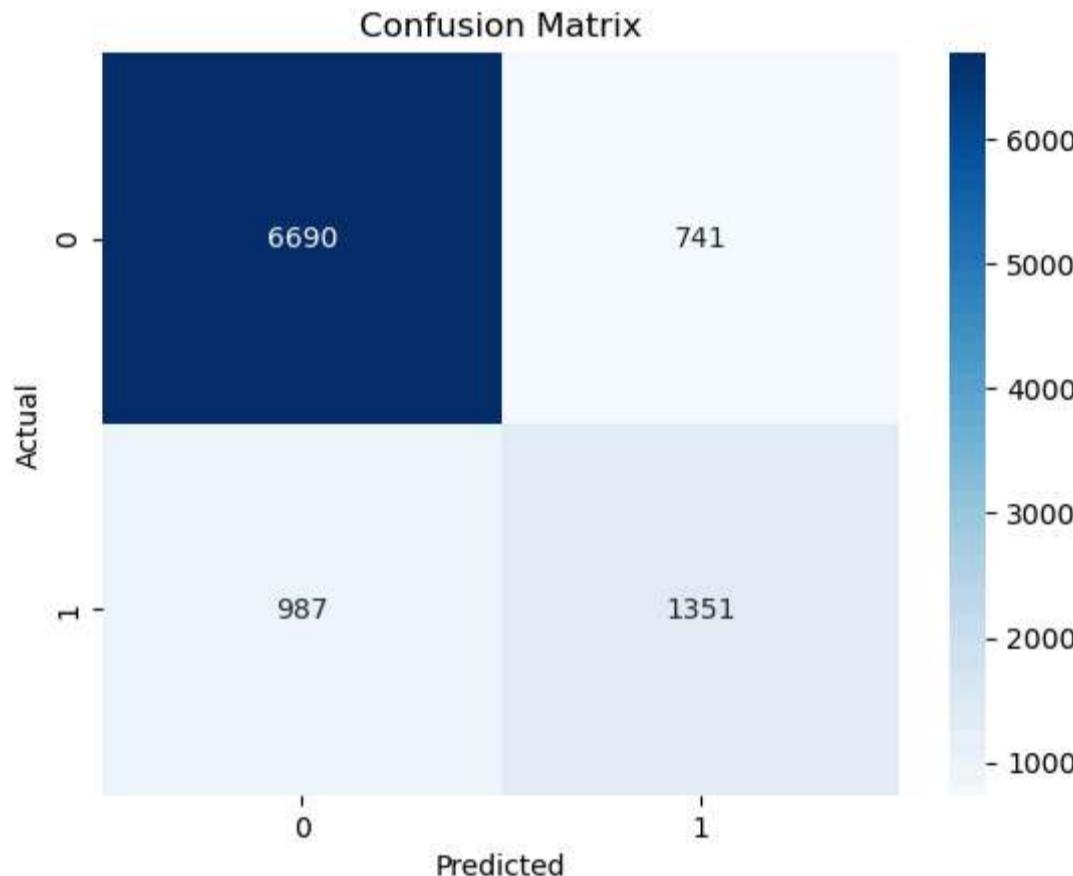
## Confusion matrix

```
In [352... confusion_matrix(y_test,y_pred)
```

```
Out[352... array([[6690, 741],
   [987, 1351]], dtype=int64)
```

```
In [353... sns.heatmap(confusion_matrix(y_test,y_pred), annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
```

```
Out[353... Text(50.72222222222214, 0.5, 'Actual')
```



```
In [355... accuracy_knn=accuracy_score(y_test,y_pred)
accuracy_knn
```

```
Out[355... 0.8231139318251612
```

```
In [356... print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.87	0.90	0.89	7431
1	0.65	0.58	0.61	2338
accuracy			0.82	9769
macro avg	0.76	0.74	0.75	9769
weighted avg	0.82	0.82	0.82	9769

## Collect accuracy scores from all models

```
In [361... Model_name={"LogisticRegression":round(accuracy_log,2),
   "RandomForestClassifier":round(accuracy_rand,2),
   "DecisionTreeClassifier":round(accuracy_tree,2),
   "KNearestNeighbors":round(accuracy_knn,2) }
Model_name
```

```
Out[361... {'LogisticRegression': 0.85,
 'RandomForestClassifier': 0.85,
 'DecisionTreeClassifier': 0.85,
 'KNearestNeighbors': 0.82}
```

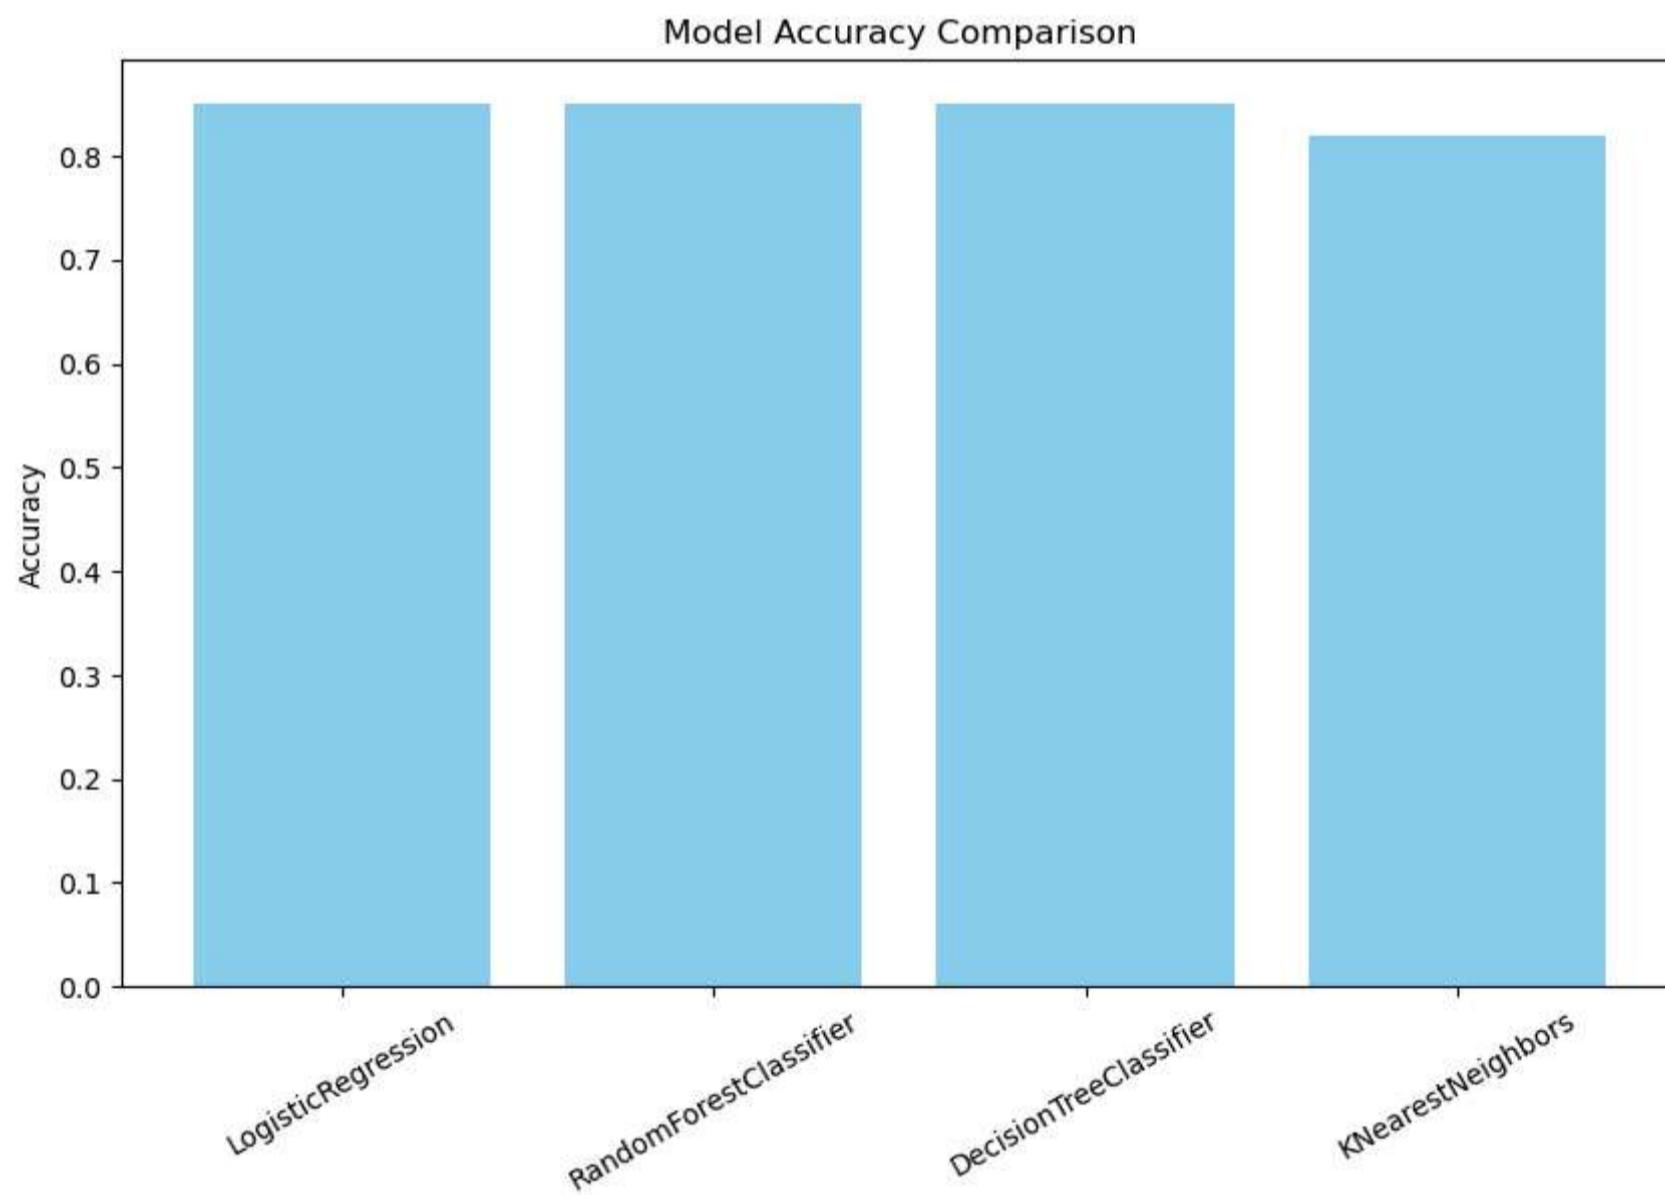
### Create a Dataframe for visualization

```
In [366... project_report=pd.DataFrame(list(Model_name.items()),columns=["Model","Accuracy"])
project_report
```

```
Out[366...      Model  Accuracy
0    LogisticRegression  0.85
1  RandomForestClassifier  0.85
2   DecisionTreeClassifier  0.85
3    KNearestNeighbors  0.82
```

```
In [368... import matplotlib.pyplot as plt
```

```
In [370... Model_name=['LogisticRegression','RandomForestClassifier','DecisionTreeClassifier','KNearestNeighbors']
Accuracies=[round(accuracy_log,2),round(accuracy_rand,2),round(accuracy_tree,2),round(accuracy_knn,2)]
plt.figure(figsize=(10,6))
plt.title('Model Accuracy Comparison')
plt.bar(Model_name,Accuracies,color='skyblue')
plt.ylabel('Accuracy')
plt.xticks(rotation=30)
plt.show()
```



The Logistic Regression, Random forest classifier and Decision tree classifier gave about 85% accuracy,while KNN gave 82% accuracy.

The Random forest as the best performing model. Because it had the highest F1 score, which balances both precision and recall.